

Assignment -1

Python Programming

Assignment Date	29 September 2022
Student Name	Vishal R
Student Roll Number	CITC1907058
Maximum Marks	2 Marks

Question 1. Download the dataset: Dataset

Importing the required libraries

```
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import numpy as np
%matplotlib inline
```

Question 2. Load the dataset.

2. Load the dataset

```
df=pd.read_csv("Churn_Modelling.csv")
```

df

```
df=pd.read_csv("/content/drive/MyDrive/Churn_Modelling.csv")
df
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCreditCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	
...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	

Question 3. Perform Below Visualizations.

● Univariate Analysis

Code:

```
df.dtypes
```

```
df['Age'].value_counts()
```

```
sns.kdeplot(df['Age'])
```

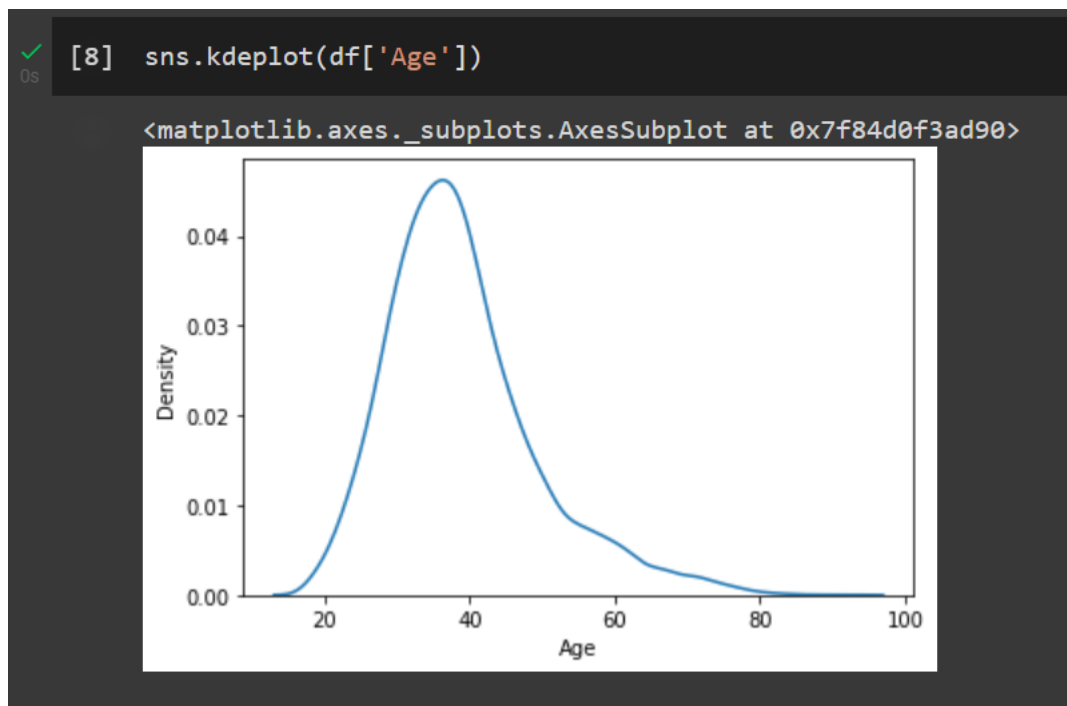
```
df.dtypes
```

RowNumber	int64
CustomerId	int64
Surname	object
CreditScore	int64
Geography	object
Gender	object
Age	int64
Tenure	int64
Balance	float64
NumOfProducts	int64
HasCrCard	int64
IsActiveMember	int64
EstimatedSalary	float64
Exited	int64
dtype:	object

```
df['Age'].value_counts()
```

37	478
38	477
35	474
36	456
34	447
...	
92	2
82	1
88	1
85	1
83	1

Name: Age, Length: 70, dtype: int64



- **Bi - Variate Analysis**

Code:

#1.

```
df.corr()
```

#2.

```
import seaborn as sns
```

```
sns.heatmap(df.corr())
```

#3.

```
import statsmodels.api as sm
```

```
#define response variable
```

```
y = df['Age']
```

```
#define explanatory variable
```

```
x = df[['Exited']]
```

```
#add constant to predictor variables
```

```
x = sm.add_constant(x)
```

```
#fit linear regression model
```

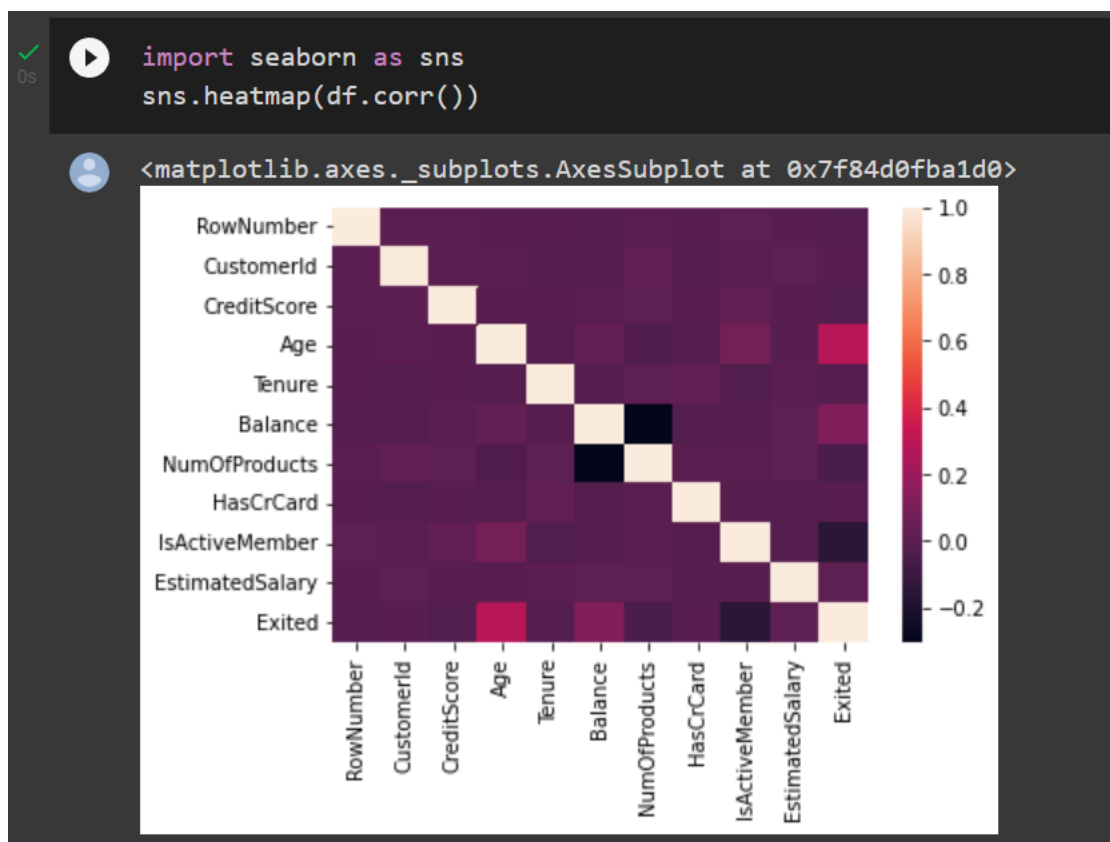
```
model = sm.OLS(y, x).fit()
```

```
#view model summary
```

```
print(model.summary())
```

0s df.corr()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber	1.000000	0.004202	0.005840	0.000783	-0.006495	-0.009067	0.007246	0.000599	0.012044	-0.005988	-0.016571
CustomerId	0.004202	1.000000	0.005308	0.009497	-0.014883	-0.012419	0.016972	-0.014025	0.001665	0.015271	-0.006248
CreditScore	0.005840	0.005308	1.000000	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Age	0.000783	0.009497	-0.003965	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	-0.006495	-0.014883	0.000842	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	-0.009067	-0.012419	0.006268	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.007246	0.016972	0.012238	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	0.000599	-0.014025	-0.005458	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.012044	0.001665	0.025651	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000		
EstimatedSalary	-0.005988	0.015271	-0.001384	-0.007201	0.007784	0.012797	0.014204	-0.009933		1.000000	
Exited	-0.016571	-0.006248	-0.027094	0.285323	-0.014001	0.118533	-0.047820	-0.007138			1.000000



```

import statsmodels.api as sm
#define response variable
y = df['Age']

#define explanatory variable
x = df[['Exited']]

#add constant to predictor variables
x = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y, x).fit()

#view model summary
print(model.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Age      R-squared:                0.081
Model:                  OLS      Adj. R-squared:            0.081
Method:                 Least Squares      F-statistic:           886.1
Date:                   Tue, 04 Oct 2022    Prob (F-statistic):     1.24e-186
Time:                   13:41:38           Log-Likelihood:        -37266.
No. Observations:      10000           AIC:                  7.454e+04
Df Residuals:          9998           BIC:                  7.455e+04
Df Model:               1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                37.4084      0.113      332.078      0.000      37.188      37.629
Exited                7.4296      0.250      29.767      0.000       6.940       7.919
=====
Omnibus:                1974.048    Durbin-Watson:           2.027
Prob(Omnibus):           0.000    Jarque-Bera (JB):        4381.188
Skew:                    1.136    Prob(JB):                 0.00
Kurtosis:                5.314    Cond. No.                 2.60
=====

```

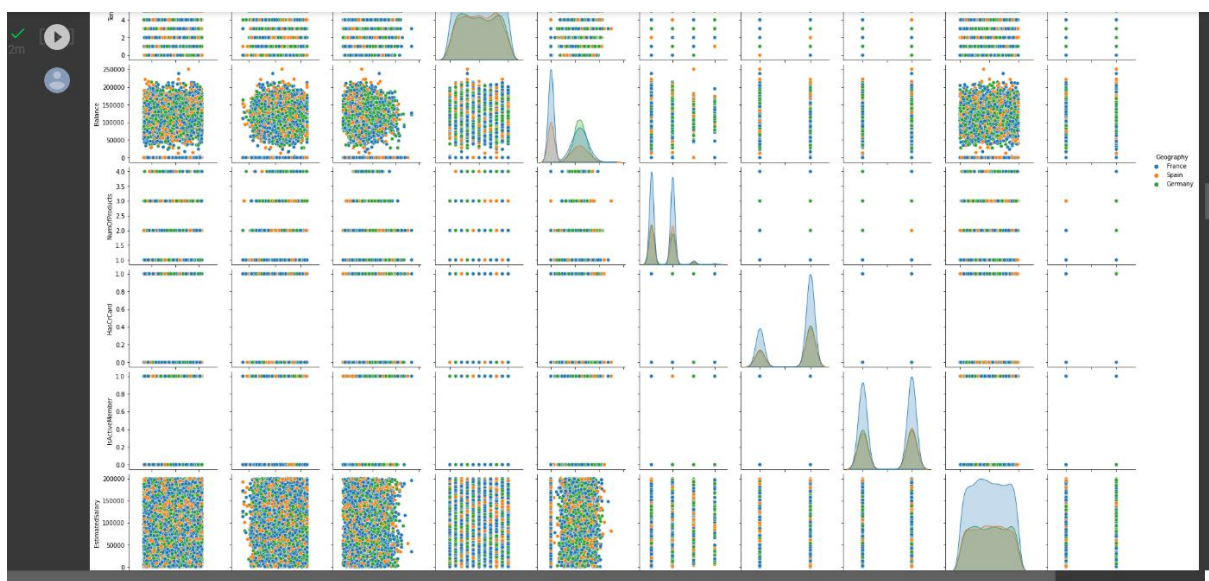
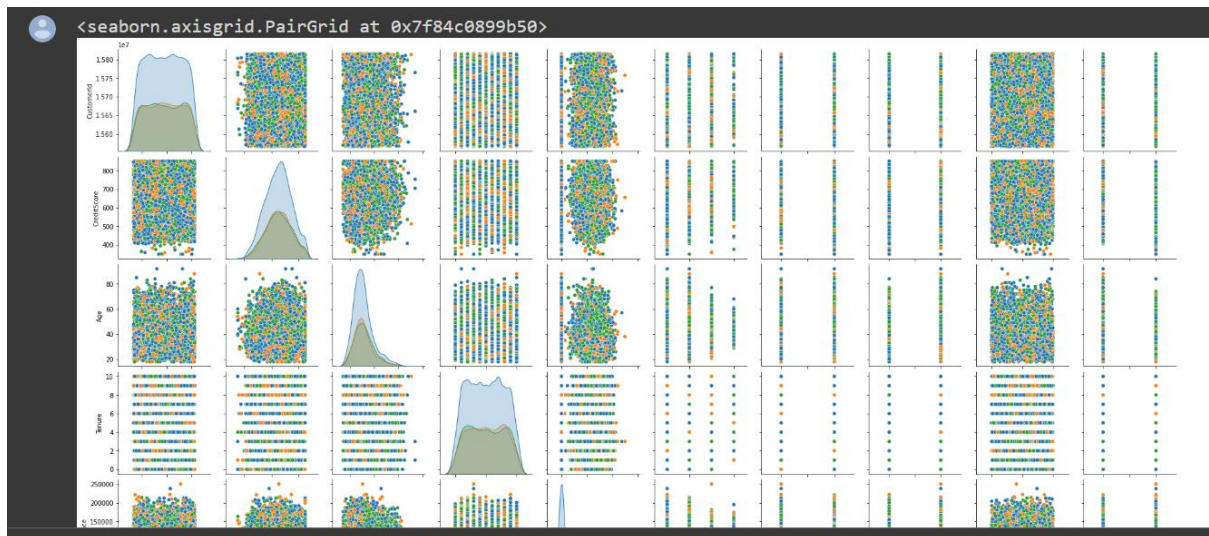
- Multi - Variate Analysis

Code

```

sns.pairplot(data=df[['CustomerId', 'CreditScore', 'Gender', 'Age', 'Tenure',
'Geography', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
'EstimatedSalary', 'Exited']], hue='Geography')

```




Question 4. Perform descriptive statistics on the dataset.

Code:


```
df.describe()
```

```
df.sum()
```

✓ 0s  `df.describe()`

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.705100
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.451100
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000



✓ 0s  `df.sum()`

RowNumber	50005000
CustomerId	156909405694
Surname	HargraveHillOnioBoniMitchellChuBartlettObinnaH...
CreditScore	6505288
Geography	FranceSpainFranceFranceSpainSpainFranceGermany...
Gender	FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...
Age	389218
Tenure	50128
Balance	764858892.88
NumOfProducts	15302
HasCrCard	7055
IsActiveMember	5151
EstimatedSalary	1000902398.81
Exited	2037
dtype:	object

Question 5. Handle the Missing values.

Code:

```
df.fillna(df.mean())
```

df.fillna(df.mean())

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame indices (only one column is present). This behavior will change in a future pandas release. Use df.dropna() instead to explicitly drop rows with any NA in the index.

"""Entry point for launching an IPython kernel.

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCreditCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	
...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	

Question 6. Find the outliers and replace the outliers

Code:

```
df["Tenure"] = np.where(df["Tenure"] > 10, np.median(df["Tenure"]),
df["Tenure"])
```

[18] df["Tenure"] = np.where(df["Tenure"] > 10, np.median(df["Tenure"]), df["Tenure"])

0	2
1	1
2	8
3	1
4	2
...	...
9995	5
9996	10
9997	7
9998	3
9999	4

Name: Tenure, Length: 10000, dtype: object

Question 7. Check for Categorical columns and perform encoding.

Code:

```
x=list(df.columns)
for i in x:
    print(pd.Categorical(df[i]))
    print("\n\n\n")
```

```
[19] x=list(df.columns)
     for i in x:
         print(pd.Categorical(df[i]))
         print("\n\n\n")

[2, 1, 8, 1, 2, ..., 5, 10, 7, 3, 4]
Length: 10000
Categories (11, int64): [0, 1, 2, 3, ..., 7, 8, 9, 10]

[0.00, 83807.86, 159660.80, 0.00, 125510.82, ..., 0.00, 57369.61, 0.00, 75075.31, 130142.79]
Length: 10000
Categories (6382, float64): [0.00, 3768.69, 12459.19, 14262.80, ..., 221532.80, 222267.63,
238387.56, 250898.09]
```

Question 8. Split the data into dependent and independent variables.

Code:

```
dependent=df[x[:2]]
independent=df[x[2:]]
print("dependent variables\n",dependent.head())
print("\n\nindependent variables\n",independent.head())
```

```
[20] dependent=df[x[:2]]
      independent=df[x[2:]]

print("dependent variables\n",dependent.head())
print("\n\nindependent variables\n",independent.head())
```

dependent variables

	RowNumber	CustomerId
0	1	15634602
1	2	15647311
2	3	15619304
3	4	15701354
4	5	15737888

independent variables

	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	\
0	Hargrave	619	France	Female	42	2	0.00	
1	Hill	608	Spain	Female	41	1	83807.86	
2	Onio	502	France	Female	42	8	159660.80	
3	Boni	699	France	Female	39	1	0.00	
4	Mitchell	850	Spain	Female	43	2	125510.82	

Question 9. Scale the independent variables

Code:

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["RowNumber"]] = scaler.fit_transform(df[["RowNumber"]])
print(df.head())
```

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["RowNumber"]] = scaler.fit_transform(df[["RowNumber"]])
print(df.head())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	\
0	0.0000	15634602	Hargrave	619	France	Female	42	2	
1	0.0001	15647311	Hill	608	Spain	Female	41	1	
2	0.0002	15619304	Onio	502	France	Female	42	8	
3	0.0003	15701354	Boni	699	France	Female	39	1	
4	0.0004	15737888	Mitchell	850	Spain	Female	43	2	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	\
0	0.00	1	1	1	101348.88	
1	83807.86	1	0	1	112542.58	
2	159660.80	3	1	0	113931.57	
3	0.00	2	0	0	93826.63	
4	125510.82	1	1	1	79084.10	

Exited

Question 10. Split the data into training and testing

Code:

```
from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```



```
from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```



```
(8000, 13)
(8000,)
(1000, 13)
(1000,)
(1000, 13)
(1000,)
(None, None)
```