

## Assignment -2

### Visualization and preprocessing

Assignment Date	29 September 2022
Student Name	Clifford Lebo J
Student Roll Number	CITC1907005
Maximum Marks	2 Marks

#### Question 1. Download the dataset: Dataset

##### 1. Importing required libraries

```
In [1]: import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import numpy as np
%matplotlib inline
```

#### Question 2. Load the dataset.

# 2. Load the dataset

```
df=pd.read_csv("Churn_Modelling.csv")
```

df

```
In [2]: df=pd.read_csv("Churn_Modelling.csv")
df
```

```
Out[2]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSa
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows × 14 columns

### Question 3. Perform Below Visualizations.

- **Univariate Analysis**

Code:

```
df.dtypes
```

```
df['Age'].value_counts()
```

```
sns.kdeplot(df['Age'])
```

### 3.1.univariate Analysis

```
In [3]: df.dtypes
```


```
Out[3]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object
```

```
In [4]: df['Age'].value_counts()
```

```
Out[4]: 37    478
38    477
35    474
36    456
34    447
...
92      2
82      1
88      1
85      1
83      1
Name: Age, Length: 70, dtype: int64
```

```
sns.kdeplot(df['Age'])
```

```
<AxesSubplot:xlabel='Age', ylabel='Density'>
```



- **Bi - Variate Analysis**

Code:

#1.

```
df.corr()
```

#2.

```
import seaborn as sns
```

```
sns.heatmap(df.corr())
```

#3.

```
import statsmodels.api as sm
```

```
#define response variable
```

```
y = df['Age']
```

```
#define explanatory variable
```

```
x = df[['Exited']]
```

```
#add constant to predictor variables
```

```
x = sm.add_constant(x)
```

```
#fit linear regression model
```

```
model = sm.OLS(y, x).fit()
```

```
#view model summary
```

```
print(model.summary())
```

### 3.2.Bi-Variate Analysis


```
In [6]: df.corr()
```

```
Out[6]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber	1.000000	0.004202	0.005840	0.000783	-0.006495	-0.009067	0.007246	0.000599	0.012044	-0.005988	-0.016571
CustomerId	0.004202	1.000000	0.005308	0.009497	-0.014883	-0.012419	0.016972	-0.014025	0.001665	0.015271	-0.006248
CreditScore	0.005840	0.005308	1.000000	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Age	0.000783	0.009497	-0.003965	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	-0.006495	-0.014883	0.000842	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	-0.009067	-0.012419	0.006268	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.007246	0.016972	0.012238	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	0.000599	-0.014025	-0.005458	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.012044	0.001665	0.025651	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000	-0.011421	-0.156128
EstimatedSalary	-0.005988	0.015271	-0.001384	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	1.000000	0.012097
Exited	-0.016571	-0.006248	-0.027094	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	1.000000

```
In [7]: import seaborn as sns  
sns.heatmap(df.corr())
```

```
Out[7]: <AxesSubplot:>
```



```
In [8]: import statsmodels.api as sm
#define response variable
y = df['Age']

#define explanatory variable
x = df[['Exited']]

#add constant to predictor variables
x = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y, x).fit()

#view model summary
print(model.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Age      R-squared:                0.081
Model:                  OLS      Adj. R-squared:           0.081
Method:                 Least Squares      F-statistic:           886.1
Date:                   Sun, 02 Oct 2022    Prob (F-statistic):     1.24e-186
Time:                   22:44:02           Log-Likelihood:        -37266.
No. Observations:      10000             AIC:                  7.454e+04
Df Residuals:          9998              BIC:                  7.455e+04
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	37.4084	0.113	332.078	0.000	37.188	37.629
Exited	7.4296	0.250	29.767	0.000	6.940	7.919

```

=====
Omnibus:                 1974.048      Durbin-Watson:           2.027
Prob(Omnibus):            0.000      Jarque-Bera (JB):        4381.188
Skew:                     1.136      Prob(JB):                0.00
Kurtosis:                 5.314      Cond. No.                 2.60
=====

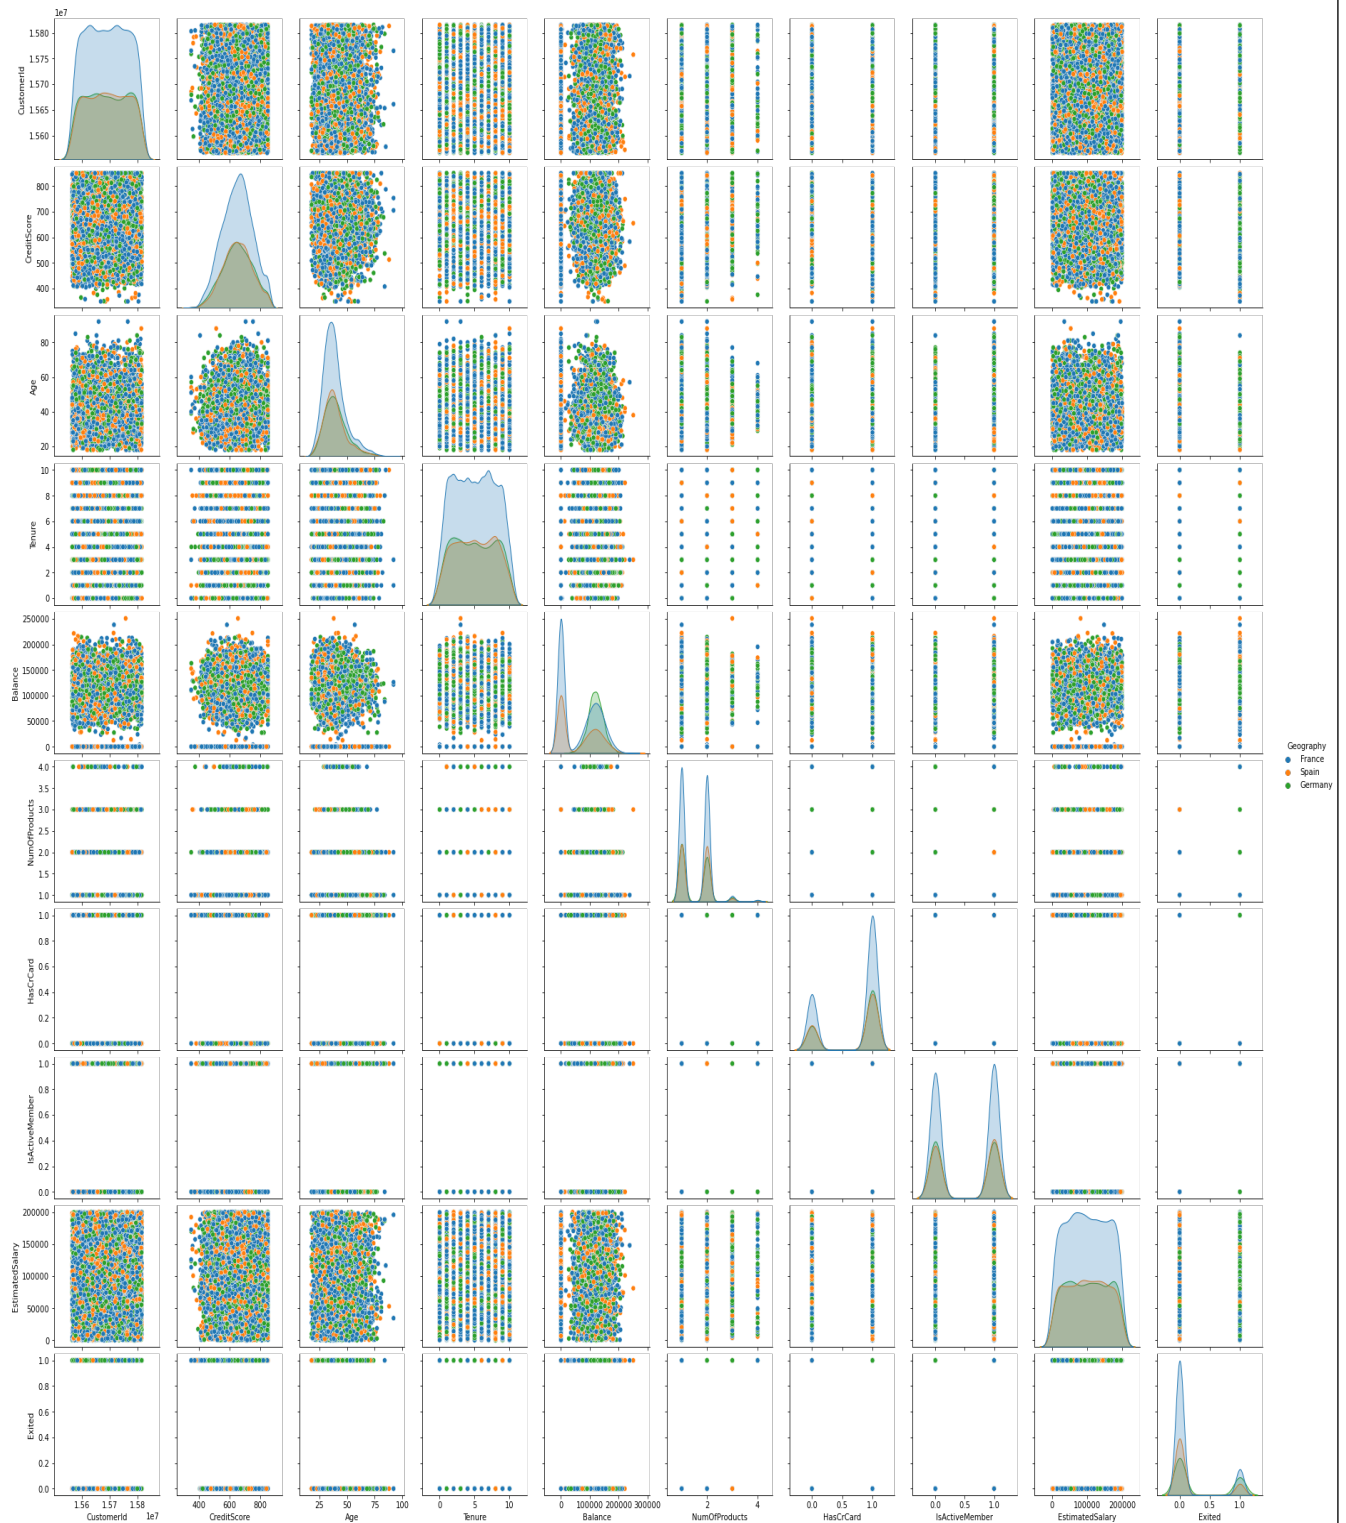
```

Notes:

## • Multi - Variate Analysis

### Code

```
sns.pairplot(data=df[['CustomerId', 'CreditScore', 'Gender', 'Age', 'Tenure',
'Geography', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
'EstimatedSalary', 'Exited']], hue='Geography')
```



**Question 4. Perform descriptive statistics on the dataset.**

Code:

```
df.describe()
```

In [11]: `df.describe()`

Out[11]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumC
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	100
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	

## Question 5. Handle the Missing values.

Code:

`df.fillna(df.mean())`

In [13]: `df.fillna(df.mean())`

C:\Users\Cliff\AppData\Local\Temp\ipykernel\_24992\634187881.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.  
`df.fillna(df.mean())`

Out[13]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.0
1	2	15647311	Hill	608	Spain	Female	41	1	83807.8
2	3	15619304	Onio	502	France	Female	42	8	159660.8
3	4	15701354	Boni	699	France	Female	39	1	0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.8
...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.0
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.6
9997	9998	15584532	Liu	709	France	Female	36	7	0.0
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.3
9999	10000	15628319	Walker	792	France	Female	28	4	130142.7

10000 rows × 10 columns

## Question 6. Find the outliers and replace the outliers



Code:

```
df["Tenure"] = np.where(df["Tenure"] > 10, np.median(df["Tenure"]),
df["Tenure"])
```

```
In [14]: df["Tenure"] = np.where(df["Tenure"] > 10, np.median(df["Tenure"]),
df["Tenure"])

Out[14]: 0      2
         1      1
         2      8
         3      1
         4      2
         ..
        9995     5
        9996    10
        9997     7
        9998     3
        9999     4
        Name: Tenure, Length: 10000, dtype: object
```

**Question 7. Check for Categorical columns and perform encoding.**

Code:

```
x=list(df.columns)
for i in x:
    print(pd.Categorical(df[i]))
    print("\n\n\n")
```

```
In [22]: x=list(df.columns)
         for i in x:
             print(pd.Categorical(df[i]))
             print("\n\n")
```

```
[1, 1, 3, 2, 1, ..., 2, 1, 1, 2, 1]
Length: 10000
Categories (4, int64): [1, 2, 3, 4]
```

```
[1, 0, 1, 0, 1, ..., 1, 1, 0, 1, 1]
Length: 10000
Categories (2, int64): [0, 1]
```

```
[1, 1, 0, 0, 1, ..., 0, 1, 1, 0, 0]
Length: 10000
Categories (2, int64): [0, 1]
```

**Question 8. Split the data into dependent and independent variables.**

Code:

```
dependent=df[x[:2]]
```

```
independent=df[x[2:]]
```

```
print("dependent variables\n",dependent.head())
```

```
print("\n\nindependent variables\n",independent.head())
```

```
In [16]: dependent=df[x[:2]]
independent=df[x[2:]]
```

```
In [17]: print("dependent variables\n",dependent.head())
print("\nindependent variables\n",independent.head())
```

dependent variables

	RowNumber	CustomerId
0	1	15634602
1	2	15647311
2	3	15619304
3	4	15701354
4	5	15737888

independent variables

	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	\
0	Hargrave	619	France	Female	42	2	0.00	
1	Hill	608	Spain	Female	41	1	83807.86	
2	Onio	502	France	Female	42	8	159660.80	
3	Boni	699	France	Female	39	1	0.00	
4	Mitchell	850	Spain	Female	43	2	125510.82	

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	1	101348.88	1
1	1	0	1	112542.58	0
2	3	1	0	113931.57	1
3	2	0	0	93826.63	0
4	1	1	1	79084.10	0

## Question 9. Scale the independent variables

Code:

```
import pandas as pd
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
df[["RowNumber"]] = scaler.fit_transform(df[["RowNumber"]])
```

```
print(df.head())
```

```
In [18]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["RowNumber"]] = scaler.fit_transform(df[["RowNumber"]])
print(df.head())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
\	0	0.0000	15634602	Hargrave	619	France	Female	42
	1	0.0001	15647311	Hill	608	Spain	Female	41
	2	0.0002	15619304	Onio	502	France	Female	42
	3	0.0003	15701354	Boni	699	France	Female	39
	4	0.0004	15737888	Mitchell	850	Spain	Female	43

	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	\
0	0.00	1	1	1	101348.88	
1	83807.86	1	0	1	112542.58	
2	159660.80	3	1	0	113931.57	
3	0.00	2	0	0	93826.63	
4	125510.82	1	1	1	79084.10	

	Exited
0	1
1	0
2	1
3	0
4	0

## Question 10. Split the data into training and testing

Code:

```
from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```

```
In [19]: from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```

```
(8000, 13)
(8000,)
(1000, 13)
(1000,)
(1000, 13)
(1000,)
```

```
Out[19]: (None, None)
```